

The background is a solid blue color. In the upper left quadrant, there is a complex, abstract geometric pattern composed of several thin, white, overlapping lines that form various polygons and angles. A single, thin white line extends diagonally from the top right towards the bottom right corner. Another thin white line is horizontal, positioned in the upper right area.

VaultScaler Labs

Pitch Deck

AGENDA

Executive Summary

Issues and Opportunities

Radix Core and Radix Studio Solutions

Market

Product Roadmap and Pricing

Competitive Landscape

Team

Why VaultScaler, Why Now?

Funding



Executive Summary

**VaultScaler Labs exists to
democratize and economize AI operations**

Radix Core

- Model **TRAINING** optimization
- **21% throughput gains** for model training.
A company with 100 GPUs will save roughly \$250,000/year.
- No CapEx. No stack changes
- Cluster administration policies for user fairness and cluster health

Radix Studio

- Model **INFERENCE** optimization
- Cost-aware LLM utilization
- No more AI vendor lock-in
- Optimized utilization of the right model, tool, GPU, policies, and execution order

Issues and Opportunities

Model Side (Radix Core)

*Existing GPU schedulers operate in open-loop mode, treating GPUs as monolithic units. This could create **40-60% resource imbalance** across memory, compute, and power dimensions. (1)*

Organizations are facing:

- Large model training bills for AWS, Azure, and other cloud providers
- Inefficient use of GPU infrastructure
- Blindness to cost leakage

Inference Side (Radix Studio)

*62% of organizations said a **lack of data governance is the main data challenge** inhibiting AI initiatives according to a KPMG survey. (2)*

Organizations are facing:

- AI solutions that lack the required governance trust to move forward with
- Lack of agility in AI vendor lock in and point solutions for managing LLM calls
- Unknown data quality, ownership, and usage

(1) Summarized findings from: Subramaniyan, S., & Wang, X. (2025). FC-GPU: Feedback Control GPU Scheduling for Real-time Embedded Systems. ACM Transactions on Embedded Computing Systems, 24(5s), Article 155. <https://doi.org/10.1145/3761812> Ohio State University, ACM Transactions on Embedded Computing Systems, 2025

(2) KPMG. (2025, October). Data governance in the age of AI. KPMG Belgium. <https://kpmg.com/be/en/home/insights/2025/10/ta-data-governance-in-the-age-of-ai.html>

Radix Core Solution

Radix Core enables organizations to see up to **21% model training efficiency gains**

Light

Install Radix agents on Kubernetes clusters. No stack changes. No changes initiating training jobs

Demonstrable

Leadership ready dashboard reports showing Radix performance compared to your existing scheduler

Adaptive

Radix leverages continual scheduling refinement AI to ensure the best results

Secure

Signed images/charts, SLSA Level 3 attestations and CycloneDX/SPDX SBOMs for complete provenance tracking. Restricted Kubernetes Pod Security Standards- read-only file systems, non-root execution and dropped capabilities

Radix Studio Solution

Complete LLM operations and sustainability

Reusable

Chain LLM calls, RAG queries, and tool integrations into production-ready workflows with Radix Multi-Mode Registry

Governed

Enforce compliance policies, content filtering, and budget controls at every step

Transparent

Real-time visibility into each pipeline step with full execution traces

Secure

Multi-Tenant Architecture for complete data isolation for each customer

Market

Total Addressable Market

by Segment

Tier	Core	Studio	TAM
Team	-	\$ 696,000	\$ 696,000
Professional	\$ 26,611,200	\$ 1,432,800	\$ 28,044,000
Enterprise	\$ 30,412,800	\$ 1,437,120	\$ 31,849,920
Hyperscale	\$ 2,736,000	0	\$ 2,736,000
Total	\$ 59,760,000	\$ 3,565,920	\$ 63,325,920

Serviceable Addressable Market

by Segment

	Moderate Projection	
Tier	Mkt %	Revenue
Core Professional	5%	\$ 1,330,560
Core Enterprise	4%	\$ 1,216,512
Studio Team	5%	\$ 34,800
Studio Professional	4%	\$ 1,121,760
Studio Enterprise	4%	\$ 1,273,997
Total		\$ 4,977,629



Product Roadmap and Pricing

Radix Core Professional, Radix Studio Team & Professional | *December 2025*

- \$28/GPU/mo. Up to 400 GPUs. Full trial for 14 days
- Studio Team: \$29/mo. for 1 user, 3 models & pipelines
- Studio Professional: \$199/mo. for 5 users, 20 models & pipelines

Radix Core and Studio Enterprise | *September 2026*

- Radix Core Enterprise: \$22/GPU/Mo. 1000 GPUs. Full trial for 14 days
- Radix Studio Enterprise: \$499/mo. for 10 users, 50 models & pipelines

Competitive Landscape

Radix Core

Professional Level – Limited competition. Organizations running under 400 GPUs aren't likely to have scheduling optimizations

Enterprise Level -Some competition with Ray and Slurm

Hyperscaler Level - Likely have their own proprietary stack.

Radix Studio

vs. TrueFoundry

- Targeted LLM operations solution is far less expensive
- True SaaS vs. their K8s-first approach that adds operational overhead

vs. Azure AI / AWS Bedrock

- Single platform vs. 4+ services to integrate and manage
- AI vendor and BYOG vs. their LLM lock-in
- Custom governance policies and auditability
- Predictable per-execution pricing



Team

Ian Green

Founder & CEO

Ian Green is an AI platform engineer and the founder of VaultScaler. He builds Kubernetes- and Terraform-driven systems that deploy, schedule, and observe GPU workloads across data centers—focused on reliability, security, and cost-aware performance. Previously a DevOps engineer at Vontier and an Electronic Warfare specialist in the New Zealand Defense Force, Ian brings deep CI/CD, AWS, and observability experience to production AI. He holds AWS Solutions Architect and Data Analytics certifications.

Ben Pruess

Founder, President & COO

Ben has been delivering enterprise grade solutions in healthcare tech for 20 years. He has rolled out numerous solutions and has directed a team of 35 while modernizing a \$120M platform. Ben is rigorous about delivering value to the customer. He calls Las Vegas home, along with his wife and three boys where he enjoys being sometimes right and always blessed.

Why VaultScaler, Why Now:

VaultScaler is positioned to deliver

The world needs more efficient model training

Radix Core is the first mover in a niche category that delivers immediate 21% throughput gains with zero capital expense or stack changes

Radix Studio delivers needed LLM inference performance, accuracy, observability and governance that organizations need to operationalize LLMs with trust

Funding

We will ask for **\$90,000** after initial market validation of 4 Professional customers

This will

- Bridge an 11-month path to profitability
- Get VaultScaler to \$1M ARR
- Allow quicker growth than bootstrapping
- Unlock the Enterprise market sooner

A series of white, thin, overlapping geometric lines on a dark blue background, creating a complex, abstract pattern on the left side of the slide.

THANK YOU